



华南理工大学

South China University of Technology

---

## The Experiment Report of *Machine Learning*

---

**SCHOOL:** SCHOOL OF SOFTWARE ENGINEERING

**SUBJECT:** SOFTWARE ENGINEERING

*Author:*  
Deng Huang

*Supervisor:*  
Mingkui Tan or Qingyao Wu

*Student ID:*  
201530211588

*Grade:*  
Undergraduate

December 15, 2017

# Linear Regression, Linear Classification and Gradient Descent

**Abstract**—Empirical risk minimization (ERM) is always a problem in machine learning. This paper focus on the comparison of four improvements of Stochastic gradient descent (SGD) - NAG, RMSProp, AdaDelta and Adam, and their performance.

## I. INTRODUCTION

- 1) Compare and understand the difference between gradient descent and stochastic gradient descent.
- 2) Compare and understand the differences and relationships between Logistic regression and linear classification.
- 3) Further understand the principles of SVM and practice on larger data.

## II. METHODS AND THEORY

### A. NAG

NAG (Nesterov accerlerated gradient) uses Momentum to predict next gradient, instead of using current  $\theta$

$$\begin{aligned} g_t &\leftarrow \nabla J(\theta_{t-1} - \gamma v_{t-1}) \\ v_t &\leftarrow \gamma v_{t-1} + \eta g_t \\ \theta_t &\leftarrow \theta_{t-1} - v_t \end{aligned}$$

### B. RMSProp

To solve the problem in AdaGrad, RMSProp was put forward by Hinton

$$\begin{aligned} g_t &\leftarrow \nabla J(\theta_{t-1}) \\ G_t &\leftarrow \gamma G_t + (1 - \gamma) g_t \odot g_t \\ \theta_t &\leftarrow \theta_{t-1} - \frac{\eta}{\sqrt{G_t + \epsilon}} \odot g_t \end{aligned}$$

### C. AdaDelta

Also solves the same problem as RMProp does but has no learning rate to be set.

$$\begin{aligned} g_t &\leftarrow \nabla J(\theta_{t-1}) \\ G_t &\leftarrow \gamma G_t + (1 - \gamma) g_t \odot g_t \\ \Delta\theta_t &\leftarrow -\frac{\sqrt{\Delta_{t-1} + \epsilon}}{\sqrt{G_t + \epsilon}} \odot g_t \\ \theta_t &\leftarrow \theta_{t-1} + \Delta\theta_t \\ \Delta_t &\leftarrow \gamma \Delta_{t-1} + (1 - \gamma) \Delta\theta_t \odot \Delta\theta_t \end{aligned}$$

### D. Adam

Adaptive estimates of lower-order moments. Has both the advantage of AdaGrad and RMSProp.

$$\begin{aligned} g_t &\leftarrow \nabla J(\theta_{t-1}) \\ m_t &\leftarrow \beta_1 m_{t-1} + (1 - \beta_1) g_t \\ G_t &\leftarrow \gamma G_t + (1 - \gamma) g_t \odot g_t \\ \alpha &\leftarrow \eta \frac{\sqrt{1 - \gamma^t}}{1 - \beta^t} \\ \theta_t &\leftarrow \theta_{t-1} - \alpha \frac{m_t}{\sqrt{G_t + \epsilon}} \end{aligned}$$

## III. EXPERIMENTS

### A. Dataset

Experiment uses a9a of LIBSVM Data, including 32561/16281(testing) samples and each sample has 123/123 (testing) features. Please download the training set and validation set.

### B. Implementation

1) *Core Implementation*: The python implementation was show as Figure ??.

```
if method == 'gd':
    grad = gradient(X_batch, y_batch, w)
    w -= learning_rate * grad

elif method == 'NAG':
    grad = gradient(X_batch, y_batch, w - gamma * v)
    v = gamma * v + learning_rate * grad
    w -= v

elif method == 'RMSProp':
    grad = gradient(X_batch, y_batch, w)
    G = gamma * G + (1 - gamma) * np.square(grad)
    w -= learning_rate * 100 * grad / np.sqrt(G + epsilon)

elif method == 'AdaDelta':
    grad = gradient(X_batch, y_batch, w)
    G = gamma * G + (1 - gamma) * np.square(grad)
    dw = -np.sqrt(delta + epsilon) / np.sqrt(G + epsilon) * grad
    w += dw
    delta = gamma * delta + (1 - gamma) * np.square(delta)

elif method == 'Adam':
    t = i + 1
    grad = gradient(X_batch, y_batch, w)
    m = beta * m + (1 - beta) * grad
    G = gamma * G + (1 - gamma) * np.square(grad)
    alpha = learning_rate * 100 * np.sqrt(1 - gamma ** t) / (1 - beta ** t)
    w -= alpha * m / np.sqrt(G + epsilon)
```

Fig. 1. python implementation

### 2) Logistic Regression and Linear SVM Initialization:

- weights: vector filled with ones in Logistic Regression and zeros in Linear SVM
- V: vector filled with zeros
- G: vector filled with zeros
- m: vector filled with zeros

### 3) Logistic Regression and Linear SVM Parameters:

- NAG: learning rate = 0.0001,  $\gamma = 0.9$
- RMSProp: learning rate = 0.01,  $\gamma = 0.9$ ,  $\epsilon = 1e-5$
- AdaDelta:  $\gamma = 0.9$ ,  $\epsilon = 1e-5$
- Adam: learning rate = 0.01,  $\gamma = 0.9$ ,  $\epsilon = 1e-5$ ,  $\beta = 0.9$

### 4) Logistic Regression Result: shown in Figure 2




Fig. 2. Logistic Regression

### 5) Linear SVM Result: shown in Figure 3




Fig. 3. Linear SVM

## IV. CONCLUSION

- 1) You have to be very careful when implementing algorithms or you may misunderstand what those equations means.
- 2) No matter how you initialize the weights, after gradient descent they will be steady.
- 3) All the models this paper implements perform similarly to the models [1] and [2] in python package sklearn.
- 4) RMSProp acts like Adam. NAG didn't fully converge because of the iteration rounds.
- 5) Despite of all these improvements, SGD still performs well.
- 6) Learning rates vary from different methods. They should be well optimized, but I simply times 100 and it works.

## REFERENCES

- [1] Logistic regression.
- [2] Linearsvc.